


DATA WORKSHOP



DW Poznań - Duże zbiory danych #9

2019-05-21

Czyli jak ogarnąć dane New York City Crimes

Agenda

01. **New York City Crimes** - dane
02. **Csv, pickle, hdf5, feathrow, pyarrow** - jak zapisywać dane i odczytywać.
03. **Dask, MODIN, CUdf**
04. Obliczenia - i porównania

New York City crimes

<https://data.cityofnewyork.us/Public-Safety/NYPD-Complaint-Data-Historic/qgea-i56i>

- Dane zawierają informacje na temat przestępczości w Nowym Jorku (New York City Police Department (NYPD)) za okres 2006 do końca 2017
- Mają 7,309,655 rekordów
- Plik .csv ma ponad 2 GB

Odczyt performance %timeit, %time

<https://colab.research.google.com/github/jakevdp/PythonDataScienceHandbook/blob/master/notebooks/01.07-Timing-and-Profiling.ipynb>

```
1 %timeit sum(range(100))
```

```
100000 loops, best of 3: 1.54 µs per loop
```

```
1 %%timeit -n 15
```

```
2 data['OFNS_DESC'].value_counts()
```

```
15 loops, best of 3: 783 ms per loop
```

Uruchomienie kilka razy i wyliczenie
średniej czasu uruchomienia

```
1 %%time
```

```
2 data.to_csv('rows_2.csv')
```

```
CPU times: user 2min 13s, sys: 3.69 s, total: 2min 16s  
Wall time: 2min 17s
```

Wyliczenie czasu uruchomienia
komórki

%PRUN

<https://colab.research.google.com/github/jakevdp/PythonDataScienceHandbook/blob/master/notebooks/01.07-Timing-and-Profilng.ipynb>

Wyliczenie czasu uruchomienia
pojedynczych wierszy

```
1 def sum_of_lists(N):
2     total = 0
3     for i in range(5):
4         L = [j ^ (j >> i) for j in range(N)]
5         total += sum(L)
6     return total
```

```
1 %prun sum_of_lists(1000000)
```

14 function calls in 0.674 seconds

Ordered by: internal time

ncalls	totttime	percall	cumtime	percall	filename:lineno(function)
5	0.607	0.121	0.607	0.121	<ipython-input-29-f105717832a2>:4(<listcomp>)
1	0.030	0.030	0.665	0.665	<ipython-input-29-f105717832a2>:1(sum_of_lists)
5	0.028	0.006	0.028	0.006	{built-in method builtins.sum}
1	0.009	0.009	0.674	0.674	<string>:1(<module>)
1	0.000	0.000	0.674	0.674	{built-in method builtins.exec}
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' objects}

Memory_profiler %memit

<https://pypi.org/project/memory-profiler/>



Wyliczenie pamięć użytą podczas
operacji

```
1 def sum_of_lists(N):
2     total = 0
3     for i in range(5):
4         L = [j ^ (j >> i) for j in range(N)]
5         total += sum(L)
6     return total
```


```
1 %prun sum_of_lists(1000000)
```

Csv - DataFrame: 10000000 rows

<https://towardsdatascience.com/the-best-format-to-save-pandas-data-414dca023e0d>



Pickle, hdf5, feathrow, PyArrow - NYC Data



Dask

Dask is a flexible library for parallel computing in Python.

Dask is composed of two parts:

1. **Dynamic task scheduling** optimized for computation. This is similar to *Airflow*, *Luigi*, *Celery*, or *Make*, but optimized for interactive computational workloads.
2. **“Big Data” collections** like parallel arrays, dataframes, and lists that extend common interfaces like *NumPy*, *Pandas*, or *Python iterators* to larger-than-memory or distributed environments. These parallel collections run on top of dynamic task schedulers.

<https://docs.dask.org/en/latest/>

```
import pandas as pd
df = pd.read_csv('2015-01-01.csv')
df.groupby(df.user_id).value.mean()
```

```
import dask.dataframe as dd
df = dd.read_csv('2015-*-*.csv')
df.groupby(df.user_id).value.mean().compute()
```


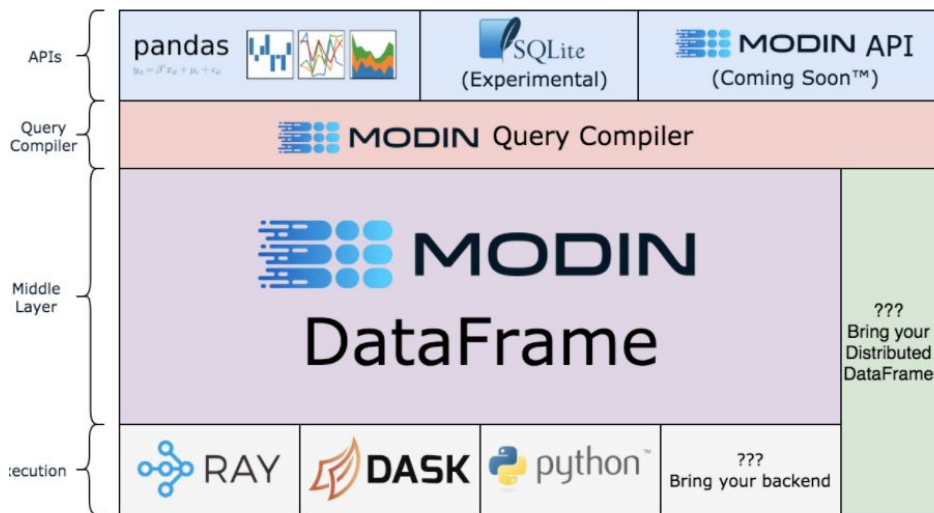
MODIN

<https://github.com/modin-project/modin>

<https://morioh.com/p/2fd9a93b6e53> - compare CUDF and MODIN

<https://towardsdatascience.com/get-faster-pandas-with-modin-even-on-your-laptops-b527a2eeda74> - COMPARE MODIN AND PANDAS

- Jest to multiprocessingowy Dataframe który może przyspieszyć wiele operacji.



CUDF


<https://rapids.ai/start.html>

<https://morioh.com/p/2fd9a93b6e53> - compare CUDF and MODIN

<https://devblogs.nvidia.com/gpu-accelerated-analytics-rapids/>

<https://towardsdatascience.com/faster-pandas-with-parallel-processing-cudf-vs-modin-f2318c594084>

- Biblioteka która cały DataFrame wysyła do pamięci GPU gdzie jest przetwarzany.
- Wszystko przez parsowanie CSV do parsowania jest robione po stronie GPU



Kolejne kroki

- Spotkanie za 2 tygodnie
- Projekt

<https://github.com/dataworkshop/dw-poznan-project>

Dziękuję